



Local installation instructions

28 November 2019

CONTEXT

HONEUR (Haematology Outcomes Network in EUROpe), a network of Haematology centers, is being established with the mutual goal of being able to track treatment outcomes in those centres and generate RWD to create valuable RWE. Each member of the network will sign a contract with Janssen agreeing to conduct studies on their RWE data and share the results.

HONEUR is a federated data network of haematology centers to support real world evidence research with a focus on outcomes management. The focus is on sharing analysis methods and results, not the granular data itself. To accommodate this objective, every data source is converted to the same data model (aka a Common Data Model - CDM). This is done locally at the participating centers, and, importantly, the data stays locally.

Since all centers are using the same model, a data analysis script can be defined centrally and sent to each of the participating centers for performing the analysis against their local data. The summary results from the individual centers can then be shared with participants through the central database, and summary analysis and visualizations can be done in the central workbench.

REQUIREMENTS

Hardware

Modern dual core processor (or better)
16 GB RAM (or more)
100 GB free disk space (or more)

Operating system

Windows 10, MacOS or Linux (Ubuntu, CentOS, Debian, ...)

Docker

- Windows: <https://docs.docker.com/docker-for-windows/install/>
- MacOS: <https://docs.docker.com/docker-for-mac/install/>
- Linux: <https://docs.docker.com/install/linux/docker-ce/ubuntu/>

Assign 2 or more CPU's, 8 GB of RAM and 100 GB of disk space to Docker in Docker Desktop.

On Linux Docker compose (v1.24 or higher) should be installed separately.

Docker images for HONEUR

A Docker Hub account with read access on the HONEUR Docker image repository (<https://hub.docker.com/u/honeur>) is required.

Please create a Docker hub account or use an existing account and request access by sending a mail to Michel Van Speybroeck (mvspeybr@its.jnj.com)

INSTALLATION COMPONENTS

The HONEUR local installation consists of the following components.

PostgreSQL database

The Postgres image for HONEUR contains the OHDSI database with 4 predefined schema's:

omopcdm

OMOP CDM (v5.3.1) schema containing the vocabulary and custom concepts for HONEUR.

webapi

Schema used by the backend of Atlas (WebAPI) to store application data

results

Schema used by the backend of Atlas (WebAPI) to store result data

scratch

Empty schema that can be freely used

See <https://www.ohdsi.org/data-standardization/the-common-data-model/> for more info about OMOP CDM

User Management

Optional component that can be used to enable security based on a local user database (as an alternative for integrating with an LDAP server).

User Management is a simple web application to create users, roles and permissions and assign permissions to roles and roles to users.

More information can be found on the HONEUR portal:

https://portal.honeur.org/group/honeur/knowledge/-/knowledge_base/user-documentation/user-management-at-the-remote-site

Atlas / WebAPI

Atlas is used to explore data in the OMOP CDM database. It's also used to import cohort definitions that are shared, run them on the local data and export the inclusion results back to the HONEUR central site.

More information can be found on the HONEUR portal

(https://portal.honeur.org/group/honeur/knowledge/-/knowledge_base/user-documentation/atlas-and-webapi-at-the-remote-site) and on the OHDSI website (<https://github.com/OHDSI/Atlas/wiki>)

Zeppelin

Zeppelin is a notebook server that is used to import and run notebooks that are shared. The results after running the analytics can be uploaded to the central side (manually or automatically). This component is also used for preparing analytical data for the distributed analytics system if your organization participates in an distributed analytics study.

More information can be found on the HONEUR portal:

https://portal.honeur.org/group/honeur/knowledge/-/knowledge_base/user-documentation/zeppelin-at-the-remote-site and on the Zeppelin website <https://zeppelin.apache.org/>

Distributed Analytics Remote

Distributed Analytics Remote is a component that is listening for distributed request messages when your organization is participating in a distributed analytics study. Requests are propagated to the computational Distributed Analytics R-Server component (see below). Results of this computation is send back to the central server where the analyst of the distributed analytics study can examine these results.

Distributed Analytics R-Server

Distributed Analytics R-Server is a computational component. This component can analyze the data that is prepared by zeppelin and send results of these computations to the distributed analytics remote component.

INSTALLATION OPTIONS

The HONEUR local components can be installed with security enabled or disabled.

When security is disabled, Atlas and Zeppelin can be accessed without login.

When security is required, the installer can choose to integrate with a local LDAP server (if available) or use the security database that comes out of the box with the secure setup.

INSTALLATION EFFORT

Installing the local HONEUR components doesn't take more than 15 minutes once a (physical or virtual) host machine - with Docker installed on it - is available.

INSTALLATION STEPS

Standard version

After the installation, 5 extra docker containers will be running:

1. Postgres v9.6.18 running on port **5444**
2. Atlas/WebAPI v2.7.1 running on Apache Tomcat 9.0.20 on port **8080**
3. Zeppelin running Jetty v9.4.14 on port **8082**
4. Distributed-Analytics R Server (no port exposed)
5. Distributed-Analytics Remote Server (no port exposed)

None of the ports that are exposed by the Docker containers (5444, 8080, 8082) should be open to the internet. The applications should only be accessible from the local network!

1. Open a terminal window (Command Prompt on Windows)

2. Download the installation script

- a. For Windows

```
curl -fsSL  
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-hon  
eur-standard.cmd --output start-honeur.cmd
```

- b. For Linux and Mac

```
curl -fsSL  
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-hon  
eur-standard.sh --output start-honeur.sh && chmod +x start-honeur.sh
```

3. Run `start-honeur.sh` (on Linux or Mac) or `start-honeur.cmd` (on Windows)
4. The script will prompt to enter the username and password of your Docker account. Make sure this docker account has read access on the HONEUR images. If you are already logged in to Docker, the script will automatically use the existing credentials.
5. Press Enter to remove the existing webapi, zeppelin and postgres container (if present).
6. After the script has downloaded the docker-compose.yml file, it will prompt you to enter a Fully Qualified Domain Name (FQDN) or IP Address of the host machine. Atlas will only be accessible on the host machine (via localhost) if you accept the default 'localhost' value.
7. The script will prompt you to enter the directory where the Zeppelin notebooks and log files should be stored.
8. The script will prompt to enter the organization you belong to (This will be given to you by the HONEUR team).

9. The Postgres database, Atlas, Zeppelin will be downloaded and will be started as Docker containers.

Secure version

After the installation, 6 extra docker containers will be running:

1. Postgres v9.6.18 running on port **5444**
2. Atlas/WebAPI v2.7.1 running on Apache Tomcat 9.0.20 on port **8080**
3. User-Management running on Apache Tomcat 9.0.33 on port **8081**
4. Zeppelin running Jetty v9.4.14 on port **8082**
5. Distributed-Analytics R Server (no port exposed)
6. Distributed-Analytics Remote Server (no port exposed)

None of the ports that are exposed by the Docker containers (5444, 8080, 8081, 8082) should be open to the internet. The applications should only be accessible from the local network!

7. Open a terminal window (Command Prompt on Windows)
8. Download the installation file

- a. For Windows

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-hon
eur-secure.cmd --output start-honeur.cmd
```

- b. For Linux and Mac

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-hon
eur-secure.sh --output start-honeur.sh && chmod +x start-honeur.sh
```

9. Run `start-honeur.sh` (on Linux or Mac) or `start-honeur.cmd` (on Windows)
10. The script will prompt you to enter the username and password of your Docker account. Make sure this Docker account has read access on the HONEUR images. If you are already logged in to docker, the script will automatically use the existing credentials.
11. Press Enter to remove existing webapi, zeppelin and postgres containers (if present).
12. After the script has downloaded the docker-compose.yml file, it will prompt you to enter a Fully Qualified Domain Name (FQDN) or IP Address of the host machine. Atlas and User Management will only be accessible on the host machine (via localhost) if you accept the default 'localhost' value.
13. The script will prompt you to enter the directory of where the Zeppelin notebooks and log files should be stored.

-
14. The script will prompt to choose between an LDAP or a JDBC connection for authentication. If you choose LDAP, the script will prompt to enter the following LDAP properties:
 - a. security.ldap.url: ldap://ldap.forumsys.com:389
 - b. security.ldap.system.username: cn=read-onlyadmin,dc=example,dc=com
 - c. security.ldap.system.password: password
 - d. security.ldap.baseDn: dc=example,dc=com
 - e. security.ldap.dn: uid={0},dc=example,dc=com
 15. The script will prompt to enter the following User Management properties:
 - a. usermgmt admin username: admin
 - b. usermgmt admin password: admin
 16. The script will prompt to enter the organization you belong to (This will be given to you by the HONEUR team).
 17. The Postgres database, Atlas, Zeppelin, User Mgmt and the distributed analytics applications will be downloaded and will be started as Docker containers

Installation Notes

Downloading the Docker images can take some time.

The following ports are used on the host machine:

- Postgres: 5444
- Atlas/WebAPI: 8080
- User management: 8081
- Zeppelin: 8082

Please contact the HONEUR support team if these ports are not available or should be changed.

The database data will be stored in a Docker volume named "pgdata"

HOW TO ACCESS THE HONEUR INSTALLATION

After the installation is completed. All HONEUR components are available and can be accessed as follows:

Postgres DB

1. Open a SQL Client (e.g.: <https://dbeaver.io/>)
2. Connect to the Postgres database running on:
Host: hostname or IP address of the host machine
Port: **5444**
Database: **OHDSI**
Username: honeur, Password: honeur (to select, update and delete data) or
Username: honeur_admin, Password: honeur_admin (to make DB changes)
Note: *Please change the default password of the honeur and honeur_admin DB accounts*
Changing the password of the honeur and honeur_admin account will not have impact on the correct working of the installation.
3. The “omopcdm”, “results”, “webapi” and scratch schemas can be accessed in the database.

Atlas

1. Open a modern browser (e.g. Google Chrome)
2. Navigate to **[http://\[hostname or IP of host machine\]:8080/atlas](http://[hostname or IP of host machine]:8080/atlas)**
3. The Atlas home page will be displayed.

User Management (optional)

1. Open a modern browser (e.g. Google Chrome)
2. Navigate to **[http://\[hostname or IP of host machine\]:8081](http://[hostname or IP of host machine]:8081)**
3. The user management home page will be displayed.

Zeppelin

1. Open a modern browser (e.g. Google Chrome)
2. Navigate to **[http://\[hostname or IP of host machine\]:8082](http://[hostname or IP of host machine]:8082)**
3. The Zeppelin home page will be displayed.

RUNNING THE ETL

The ETL code to map the source data to the OMOP Common Data Model (CDM) will be custom build for each data center. The script to execute the ETL will be provided.

POST ETL INSTALLATION STEPS

Add constraints and indexes to the OMOP CDM tables

After the ETL is successfully executed, it's recommended to add the constraints and indexes to the OMOP CDM tables. It will improve the performance and reduce the risk of corrupt data in the database.

Installation steps

1. Open a terminal window (Command Prompt on Windows)
2. Download the installation file

- a. For Windows

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-omo
pcdm-indexes-and-constraints.cmd --output
start-omopcdm-indexes-and-constraints.cmd
```

- b. For Linux and Mac

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-omo
pcdm-indexes-and-constraints.sh --output
start-omopcdm-indexes-and-constraints.sh && chmod +x
start-omopcdm-indexes-and-constraints.sh
```

1. Run `start-omopcdm-indexes-and-constraints.sh` (on Linux or Mac) or
`start-omopcdm-indexes-and-constraints.cmd` (on Windows)

The program will prompt to enter the username and password of your Docker account. Make sure this Docker account has read access on the HONEUR images. If you are already logged in to Docker, the script will automatically use the existing credentials. Press Enter to remove existing `omop-indexes-and-constraints` container (if present). After the script has downloaded the `docker-compose.yml` and `setup-conf/setup.yml` files, it will start creating the indexes and constraints on the OMOP CDM tables.

Update custom concepts in the OMOP CDM DB

When new custom concepts are available, they can be easily loaded in the OMOP CDM database.

Installation steps

1. Open a terminal window (Command Prompt on Windows)
2. Download the installation file

- For Windows: `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-omopcdm-custom-concepts-update.cmd --output start-omopcdm-custom-concepts-update.cmd`
 - For Linux and Mac: `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-omopcdm-custom-concepts-update.sh --output start-omopcdm-custom-concepts-update.sh && chmod +x start-omopcdm-custom-concepts-update.sh`
3. Run `start-omopcdm-custom-concepts-update.sh` (on Linux or Mac) or `start-omopcdm-custom-concepts-update.cmd` (on Windows)
 4. The program will prompt you for username and password for your docker account. Make sure this docker account has read access on the honeur images. If you are already logged in to docker, the program will automatically use the existing credentials.
 5. Press Enter to remove the existing `omop-cdm-custom-concepts` container.
 6. After the program has downloaded the `docker-compose.yml` and `setup-conf/setup.yml` files, it will start importing the custom concepts in the OMOP CDM database.

QA database

QA database can be used as a test database. It's an exact replica of the full database installed with the script `start-honneur.cmd` (on windows) or `start-honneur.sh` (on Linux or Mac). It is primarily used for testing scripts on the data in OMOP CDM schema.

Installation steps

1. Open a terminal window (Command Prompt on Windows)
2. Download the installation file
 - For Windows: `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-qa-database.cmd --output start-qa-database.cmd`
 - For Linux and Mac `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-qa-database.sh --output start-qa-database.sh && chmod +x start-qa-database.sh`
3. Run `start-qa-database.sh` (on Linux or Mac) or `start-qa-database.cmd` (on Windows)
4. The program will prompt you for username and password for your docker account. Make sure this docker account has read access on the honeur images. If you are already logged in to docker, the program will automatically use the existing credentials.
5. Press Enter to remove existing `postgres-qa` and `webapi-source-qa-enable` container.
6. After the program has downloaded the `docker-compose.yml` and `setup-conf/setup.yml` files, it will start initializing the QA database and insert the source of the QA database inside the original database.

7. Restart webapi/atlas container to make the source available in the webapi/atlas instance. Do this using the following command: `docker restart webapi`

Removal steps

1. Open a terminal window (Command Prompt on Windows)
2. Download the installation file
 - For Windows: `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/remove-qa-database.cmd --output remove-qa-database.cmd`
 - For Linux and Mac: `curl -fsSL https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/remove-qa-database.sh --output remove-qa-database.sh && chmod +x remove-qa-database.sh`
3. Run `remove-qa-database.sh` (on Linux or Mac) or `remove-qa-database.cmd` (on Windows)
4. The program will prompt you for username and password for your docker account. Make sure this docker account has read access on the honeur images. If you are already logged in to docker, the program will automatically use the existing credentials.
5. Press Enter to remove existing postgres-qa and webapi-source-qa-disable container.
6. After the program has downloaded the docker-compose.yml and setup-conf/setup.yml files, it will start removing the QA database and removing the source of the QA database inside the original database.

MAINTENANCE

Starting and stopping the Docker Containers

The Docker containers for HONEUR will automatically restart when the host machine is restarted. To manually start and stop the containers, open a terminal window and run one of the following commands.

Postgres:

- **to start:** `docker start postgres`
- **to stop:** `docker stop postgres`

Atlas

- **to start:** `docker start webapi`
- **to stop:** `docker stop webapi`

User Management

- **to start:** `docker start user-mgmt`

- **to stop:** `docker stop user-mgmt`

Zeppelin

- **to start:** `docker start zeppelin`
- **to stop:** `docker stop zeppelin`

Distributed Analytics Remote

- **to start:** `docker start distributed-analytics-remote`
- **to stop:** `docker stop distributed-analytics-remote`

Distributed Analytics R-Server

- **to start:** `docker start distributed-analytics-r-server`
- **to stop:** `docker stop distributed-analytics-r-server`

Backup and restore of the database

Backup

Download and run the script 'backup-database.sh':

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/backup-database
.sh --output backup-database.sh && chmod +x backup-database.sh
```

The backup script will create a tar file name 'postgres_backup_<date_time>.tar.bz2' in the current directory. Creating the backup file can take a long time depending on the size of the database.

Copy the backup file to a save location for long term storage.

Restore

Download the script 'restore-database.sh'

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/restore-databas
e.sh --output restore-database.sh && chmod +x restore-database.sh
```

Run the script and provide the name of the backup file as parameter. The name of the target volume can be provided as a second argument. The backup will be restored in the pgdata volume if the target volume is not provided. The backup file should be present in the folder where the script is executed.

```
./restore-database postgres_backup_<date_time>.tar.bz2
```

Hot snapshot of the database

The database volume can be copied to a new volume (with a different name) to take a snapshot of the current database state. Download the script 'clone-volume.sh'

```
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/clone-volume.sh
--output clone-volume.sh && chmod +x clone-volume.sh
```

Run the script, provide the source volume as first parameter and the target volume as second parameter.

```
./clone-volume.sh pgdata pgdata_snapshot1
```

Accessing log files

Console logs

Docker provides real time access to the applications logs.

- Atlas: `docker logs webapi -f`
- Zeppelin: `docker logs zeppelin -f`
- Postgres: `docker logs postgres -f`

Logs files

The log files of Zeppelin are available on the host machine where the Docker containers are running. The location of the logs files can be configured during the setup. The default location is `/zeppelin/logs`

The logs files of Atlas are available within the Docker container and can be copied to the host machine if needed. The logs files are available in the `/usr/local/tomcat/logs` directory.

Execute the following commands to inspect the log files:

- `docker exec -it webapi /bin/bash`
- `cd logs`
- `tail -200 webapi.log` or `tail catalina.<yyyy-mm-dd>.log`

Execute the following command to copy all log files to a logs directory below the current directory on host machine:

- `docker cp webapi:/usr/local/tomcat/logs .`

FURTHER READING

More documentation is available in the user documentation on the HONEUR portal:

<https://portal.honeur.org/group/honeur/knowledge>

ADDENDUM: SECURITY SCAN OF DOCKER CONTAINERS

The following images were scanned for vulnerabilities:

- honeur/postgres
- honeur/webapi-atlas
- honeur/user-mgmt
- honeur/zeppelin
- honeur-distributed-analytics

Vulnerability scanner

Docker bench security is used for the vulnerability scan on the images and containers.

<https://github.com/docker/docker-bench-security>

The Docker Bench for Security is a script that checks for dozens of common best-practices around deploying Docker containers in production. The tests are all automated and are inspired by the CIS Docker Community Edition Benchmark v1.2.0.

Results Standard Setup

Command

```
export DOCKER_CONTENT_TRUST=1
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-honeur-standard.sh --output start-honeur.sh && chmod +x start-honeur.sh
./start-honeur.sh
wget https://github.com/docker/docker-bench-security/archive/v1.3.5.tar.gz
tar -xzf v1.3.5.tar.gz
cd docker-bench-security-1.3.5
sudo sh docker-bench-security.sh
```

Output

```
# -----
# Docker Bench for Security v1.3.5
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker Benchmark v1.2.0.
# -----
```

Initializing Mon May 4 11:03:56 UTC 2020

[INFO] 4 - Container Images and Build File

```
[WARN] 4.1 - Ensure a user for the container has been created
[WARN] * Running as root: zeppelin
[WARN] * Running as root: distributed-analytics-r-server
[NOTE] 4.2 - Ensure that containers use only trusted base images
[NOTE] 4.3 - Ensure that unnecessary packages are not installed in the container
[NOTE] 4.4 - Ensure images are scanned and rebuilt to include security patches
[WARN] 4.5 - Ensure Content trust for Docker is Enabled
[PASS] 4.6 - Ensure that HEALTHCHECK instructions have been added to container images
[INFO] 4.7 - Ensure update instructions are not use alone in the Dockerfile
[INFO] * Update instruction found: [honeur/distributed-analytics:remote-1.5]
[INFO] * Update instruction found: [honeur/distributed-analytics:r-server-1.5]
[INFO] * Update instruction found: [honeur/webapi-atlas:2.7.1-1.5-standard]
[INFO] * Update instruction found: [honeur/zeppelin:0.8.2-1.5]
[INFO] * Update instruction found: [honeur/postgres:HONEUR-9.6-omopcdm-5.3.1-webapi-2.7.1-1.5]
[NOTE] 4.8 - Ensure setuid and setgid permissions are removed
[PASS] 4.9 - Ensure that COPY is used instead of ADD in Dockerfiles
[NOTE] 4.10 - Ensure secrets are not stored in Dockerfiles
[NOTE] 4.11 - Ensure only verified packages are installed
```

[INFO] 5 - Container Runtime

```
[PASS] 5.1 - Ensure that, if applicable, an AppArmor Profile is enabled
[PASS] 5.2 - Ensure that, if applicable, SELinux security options are set
[PASS] 5.3 - Ensure Linux Kernel Capabilities are restricted within containers
[PASS] 5.4 - Ensure that privileged containers are not used
[PASS] 5.5 - Ensure sensitive host system directories are not mounted on containers
[PASS] 5.6 - Ensure sshd is not run within containers
[PASS] 5.7 - Ensure privileged ports are not mapped within containers
[NOTE] 5.8 - Ensure that only needed ports are open on the container
[PASS] 5.9 - Ensure the host's network namespace is not shared
[WARN] 5.10 - Ensure that the memory usage for containers is limited
[WARN] * Container running without memory restrictions: distributed-analytics-remote
[WARN] * Container running without memory restrictions: zeppelin
[WARN] * Container running without memory restrictions: webapi
[WARN] * Container running without memory restrictions: postgres
[WARN] * Container running without memory restrictions: distributed-analytics-r-server
```

[WARN] 5.11 - Ensure CPU priority is set appropriately on the container

[WARN] * Container running without CPU restrictions: distributed-analytics-remote

[WARN] * Container running without CPU restrictions: zeppelin

[WARN] * Container running without CPU restrictions: webapi

[WARN] * Container running without CPU restrictions: postgres

[WARN] * Container running without CPU restrictions: distributed-analytics-r-server

[WARN] 5.12 - Ensure that the container's root filesystem is mounted as read only

[WARN] * Container running with root FS mounted R/W: distributed-analytics-remote

[WARN] * Container running with root FS mounted R/W: zeppelin

[WARN] * Container running with root FS mounted R/W: webapi

[WARN] * Container running with root FS mounted R/W: postgres

[WARN] * Container running with root FS mounted R/W: distributed-analytics-r-server

[WARN] 5.13 - Ensure that incoming container traffic is bound to a specific host interface

[WARN] * Port being bound to wildcard IP: 0.0.0.0 in zeppelin

[WARN] * Port being bound to wildcard IP: 0.0.0.0 in zeppelin

[WARN] * Port being bound to wildcard IP: 0.0.0.0 in webapi

[WARN] * Port being bound to wildcard IP: 0.0.0.0 in postgres

[WARN] 5.14 - Ensure that the 'on-failure' container restart policy is set to '5'

[WARN] * MaximumRetryCount is not set to 5: distributed-analytics-remote

[WARN] * MaximumRetryCount is not set to 5: zeppelin

[WARN] * MaximumRetryCount is not set to 5: webapi

[WARN] * MaximumRetryCount is not set to 5: postgres

[WARN] * MaximumRetryCount is not set to 5: distributed-analytics-r-server

[PASS] 5.15 - Ensure the host's process namespace is not shared

[PASS] 5.16 - Ensure the host's IPC namespace is not shared

[PASS] 5.17 - Ensure that host devices are not directly exposed to containers

[PASS] 5.18 - Ensure that the default ulimit is overwritten at runtime if needed

[PASS] 5.19 - Ensure mount propagation mode is not set to shared

[PASS] 5.20 - Ensure the host's UTS namespace is not shared

[PASS] 5.21 - Ensure the default seccomp profile is not Disabled

[NOTE] 5.22 - Ensure docker exec commands are not used with privileged option

[NOTE] 5.23 - Ensure that docker exec commands are not used with the user=root option

[PASS] 5.24 - Ensure that cgroup usage is confirmed

[PASS] 5.25 - Ensure that the container is restricted from acquiring additional privileges

[PASS] 5.26 - Ensure that container health is checked at runtime

[INFO] 5.27 - Ensure that Docker commands always make use of the latest version of their image

[WARN] 5.28 - Ensure that the PIDs cgroup limit is used

[WARN] * PIDs limit not set: distributed-analytics-remote

[WARN] * PIDs limit not set: zeppelin

[WARN] * PIDs limit not set: webapi

[WARN] * PIDs limit not set: postgres

[WARN] * PIDs limit not set: distributed-analytics-r-server

[PASS] 5.29 - Ensure that Docker's default bridge 'docker0' is not used

[PASS] 5.30 - Ensure that the host's user namespaces are not shared

[PASS] 5.31 - Ensure that the Docker socket is not mounted inside any containers

Results Secure Setup

Command

```
export DOCKER_CONTENT_TRUST=1
curl -fsSL
https://github.com/solventrix/Honeur-Setup/releases/download/v1.5/start-honeur-se
cure.sh --output start-honeur.sh && chmod +x start-honeur.sh
./start-honeur.sh
wget https://github.com/docker/docker-bench-security/archive/v1.3.5.tar.gz
tar -xzf v1.3.5.tar.gz
cd docker-bench-security-1.3.5
sudo sh docker-bench-security.sh
```

Output

```
# -----
# Docker Bench for Security v1.3.5
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker Benchmark v1.2.0.
# -----
```

Initializing Tue 04 Feb 2020 02:03:30 PM CET

[INFO] 4 - Container Images and Build File

```
[WARN] 4.1 - Ensure a user for the container has been created
[WARN] * Running as root: zeppelin
[WARN] * Running as root: distributed-analytics-r-server
[NOTE] 4.2 - Ensure that containers use only trusted base images
[NOTE] 4.3 - Ensure that unnecessary packages are not installed in the container
[NOTE] 4.4 - Ensure images are scanned and rebuilt to include security patches
[WARN] 4.5 - Ensure Content trust for Docker is Enabled
[PASS] 4.6 - Ensure that HEALTHCHECK instructions have been added to container images
[INFO] 4.7 - Ensure update instructions are not use alone in the Dockerfile
[INFO] * Update instruction found: [honeur/user-mgmt:1.5]
[INFO] * Update instruction found: [honeur/distributed-analytics:remote-1.5]
[INFO] * Update instruction found: [honeur/distributed-analytics:r-server-1.5]
[INFO] * Update instruction found: [honeur/webapi-atlas:2.7.1-1.5-secure]
[INFO] * Update instruction found: [honeur/zeppelin:0.8.2-1.5]
[INFO] * Update instruction found: [honeur/postgres:HONEUR-9.6-omopcdm-5.3.1-webapi-2.7.1-1.5]
[NOTE] 4.8 - Ensure setuid and setgid permissions are removed
[PASS] 4.9 - Ensure that COPY is used instead of ADD in Dockerfiles
[NOTE] 4.10 - Ensure secrets are not stored in Dockerfiles
[NOTE] 4.11 - Ensure only verified packages are installed
```

[INFO] 5 - Container Runtime

```
[PASS] 5.1 - Ensure that, if applicable, an AppArmor Profile is enabled
[PASS] 5.2 - Ensure that, if applicable, SELinux security options are set
[PASS] 5.3 - Ensure Linux Kernel Capabilities are restricted within containers
[PASS] 5.4 - Ensure that privileged containers are not used
[PASS] 5.5 - Ensure sensitive host system directories are not mounted on containers
[PASS] 5.6 - Ensure sshd is not run within containers
[PASS] 5.7 - Ensure privileged ports are not mapped within containers
[NOTE] 5.8 - Ensure that only needed ports are open on the container
[PASS] 5.9 - Ensure the host's network namespace is not shared
[WARN] 5.10 - Ensure that the memory usage for containers is limited
[WARN] * Container running without memory restrictions: distributed-analytics-remote
[WARN] * Container running without memory restrictions: webapi
[WARN] * Container running without memory restrictions: zeppelin
[WARN] * Container running without memory restrictions: user-mgmt
```

```

[WARN] * Container running without memory restrictions: postgres
[WARN] * Container running without memory restrictions: distributed-analytics-r-server
[WARN] 5.11 - Ensure CPU priority is set appropriately on the container
[WARN] * Container running without CPU restrictions: distributed-analytics-remote
[WARN] * Container running without CPU restrictions: webapi
[WARN] * Container running without CPU restrictions: zeppelin
[WARN] * Container running without CPU restrictions: user-mgmt
[WARN] * Container running without CPU restrictions: postgres
[WARN] * Container running without CPU restrictions: distributed-analytics-r-server
[WARN] 5.12 - Ensure that the container's root filesystem is mounted as read only
[WARN] * Container running with root FS mounted R/W: distributed-analytics-remote
[WARN] * Container running with root FS mounted R/W: webapi
[WARN] * Container running with root FS mounted R/W: zeppelin
[WARN] * Container running with root FS mounted R/W: user-mgmt
[WARN] * Container running with root FS mounted R/W: postgres
[WARN] * Container running with root FS mounted R/W: distributed-analytics-r-server
[WARN] 5.13 - Ensure that incoming container traffic is bound to a specific host interface
[WARN] * Port being bound to wildcard IP: 0.0.0.0 in webapi
[WARN] * Port being bound to wildcard IP: 0.0.0.0 in zeppelin
[WARN] * Port being bound to wildcard IP: 0.0.0.0 in zeppelin
[WARN] * Port being bound to wildcard IP: 0.0.0.0 in user-mgmt
[WARN] * Port being bound to wildcard IP: 0.0.0.0 in postgres
[WARN] 5.14 - Ensure that the 'on-failure' container restart policy is set to '5'
[WARN] * MaximumRetryCount is not set to 5: distributed-analytics-remote
[WARN] * MaximumRetryCount is not set to 5: webapi
[WARN] * MaximumRetryCount is not set to 5: zeppelin
[WARN] * MaximumRetryCount is not set to 5: user-mgmt
[WARN] * MaximumRetryCount is not set to 5: postgres
[WARN] * MaximumRetryCount is not set to 5: distributed-analytics-r-server
[PASS] 5.15 - Ensure the host's process namespace is not shared
[PASS] 5.16 - Ensure the host's IPC namespace is not shared
[PASS] 5.17 - Ensure that host devices are not directly exposed to containers
[PASS] 5.18 - Ensure that the default ulimit is overwritten at runtime if needed
[PASS] 5.19 - Ensure mount propagation mode is not set to shared
[PASS] 5.20 - Ensure the host's UTS namespace is not shared
[PASS] 5.21 - Ensure the default seccomp profile is not Disabled
[NOTE] 5.22 - Ensure docker exec commands are not used with privileged option
[NOTE] 5.23 - Ensure that docker exec commands are not used with the user=root option
[PASS] 5.24 - Ensure that cgroup usage is confirmed
[PASS] 5.25 - Ensure that the container is restricted from acquiring additional privileges
[PASS] 5.26 - Ensure that container health is checked at runtime
[INFO] 5.27 - Ensure that Docker commands always make use of the latest version of their image
[WARN] 5.28 - Ensure that the PIDs cgroup limit is used
[WARN] * PIDs limit not set: distributed-analytics-remote
[WARN] * PIDs limit not set: webapi
[WARN] * PIDs limit not set: zeppelin
[WARN] * PIDs limit not set: user-mgmt
[WARN] * PIDs limit not set: postgres
[WARN] * PIDs limit not set: distributed-analytics-r-server
[PASS] 5.29 - Ensure that Docker's default bridge 'docker0' is not used
[PASS] 5.30 - Ensure that the host's user namespaces are not shared
[PASS] 5.31 - Ensure that the Docker socket is not mounted inside any containers

```

Clarifications

WARNINGS

[WARN] 4.1 - Ensure a user for the container has been created

The final command or entrypoint of an image should never be run as a root user. If the container should get compromised, root access is granted to the attacker which is never good.

In the following images we prevented this:

- honeur/postgres:HONEUR-9.6-omopcdm-5.3.1-webapi-2.7.1-1.5
- honeur/webapi-atlas:2.7.1-1.5-standard
- honeur/webapi-atlas:2.7.1-1.5-secure
- honeur/user-mgmt:1.5

- honeur-distributed-analytics:remote-1.5

honeur-distributed-analytics:r-server-1.5 start as root to change the permissions of the shared volume between host, r-server and zeppelin. Once the permissions of the shared volume are changed, the r-server is run by the **nobody** user.

For Zeppelin we used apache:zeppelin:0.8.2 as the base image and only did configuration changes. Because the base image has a few programs depending on root access in the container, we decided to keep the root user.

[WARN] 4.5 - Ensure Content trust for Docker is Enabled

You can enable content trust in docker setting the environment variable DOCKER_CONTENT_TRUST to 1:

```
export DOCKER_CONTENT_TRUST=1
```

This will tell docker that it only should trust signed images.

All the images delivered by HONEUR are signed.

[WARN] 5.10 - Ensure that the memory usage for containers is limited

This are runtime warnings and are not related to the HONEUR images. CGroups can be created on the host machine and referenced in the docker containers to avoid this warning.

[WARN] 5.11 - Ensure CPU priority is set appropriately on the container

This are runtime warnings and are not related to the HONEUR images. CGroups can be created on the host machine and referenced in the docker containers to avoid this warning.

[WARN] 5.12 - Ensure that the container's root filesystem is mounted as read only

Some services on operating systems still work with access to the root filesystem. In our case the root filesystem was used to write process locks to. If we enabled read only for the root filesystem, we could not get the main programs run in the containers.

[WARN] 5.13 - Ensure that incoming container traffic is bound to a specific host interface

0.0.0.0 is a wildcard IP Address. We opted for this instead of the host IP address because our remote system will mostly run on computers with dynamic IP Addresses. It is still possible to change this for servers with a fixed IP address.

[WARN] 5.14 - Ensure that the 'on-failure' container restart policy is set to '5'

The restart policy we use is to always restart, not only on failure so MaximumRetryCount does not apply to this policy.

[WARN] 5.28 - Ensure that the PIDs cgroup limit is used

This are runtime warnings and are not directly related to the HONEUR images. CGroups can be created on host machine and referenced in the docker container to avoid this warning.